

</ Distributed transaction using Saga pattern

/>

} /> [

Speaker: Trung Tran

</ Table of contents

{01}

Context & problems

{02}

Saga - The solution

{03}

Techniques in Saga

{04}

Saga - Choreography

{05}

Saga - Orchestrator

{06}

Issues & considerations

{07}

Practical usage

{08}

Discussion & QnA

</ Prerequisites

Must have

- Basic knowledge in software architecture
- Understanding of remote communication between servers/software components
- Understanding of concurrency concepts

Nice to have

- C#, .NET, Docker knowledge
- Understanding of async communication
- Basic knowledges of distributed software system



</ Resources

Repository URL:

<https://github.com/trannamtrung1st/SagaSeminar>

Requirements:

- .NET SDK 6.0
- Docker Desktop




</ Saigon Technology

+300 software engineers

+10 years in business

3 development centers (2 offices in
Singapore and USA)

Website: <https://saigontechnology.com/>



/>

</ Saigon Technology



1 0 1 1 0 1 1 1 0 1 1 1 1 1 0 1

</ Contact

trannamtrung1st@gmail.com

<https://www.linkedin.com/in/trung-tran99/>

} /> [

/>

</>

Context & problems

01

} /> [

1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1 0 1

</ Relational database & transactions

- Transaction is a single unit of logic or work
- ACID properties:
 - Atomic
 - Consistent
 - Isolated
 - Durable

-> How about a transaction in a distributed system?

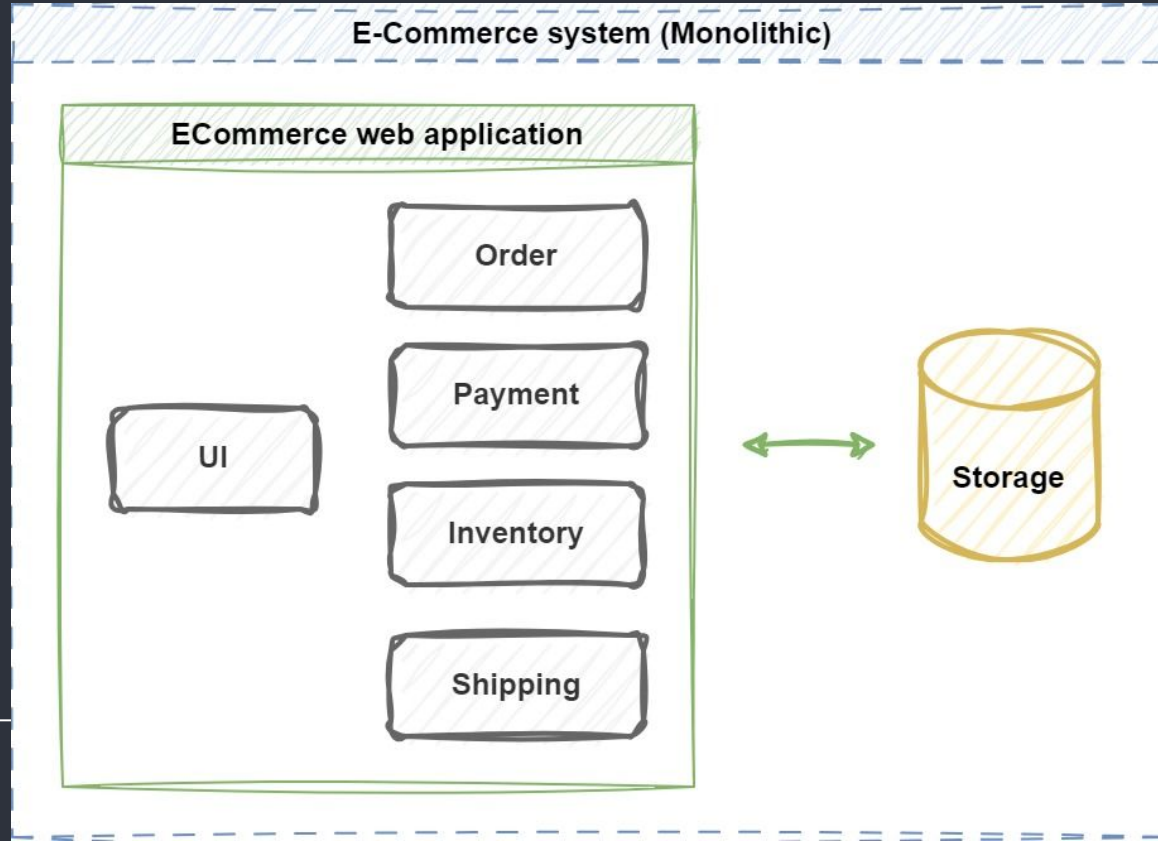


</ Distributed system

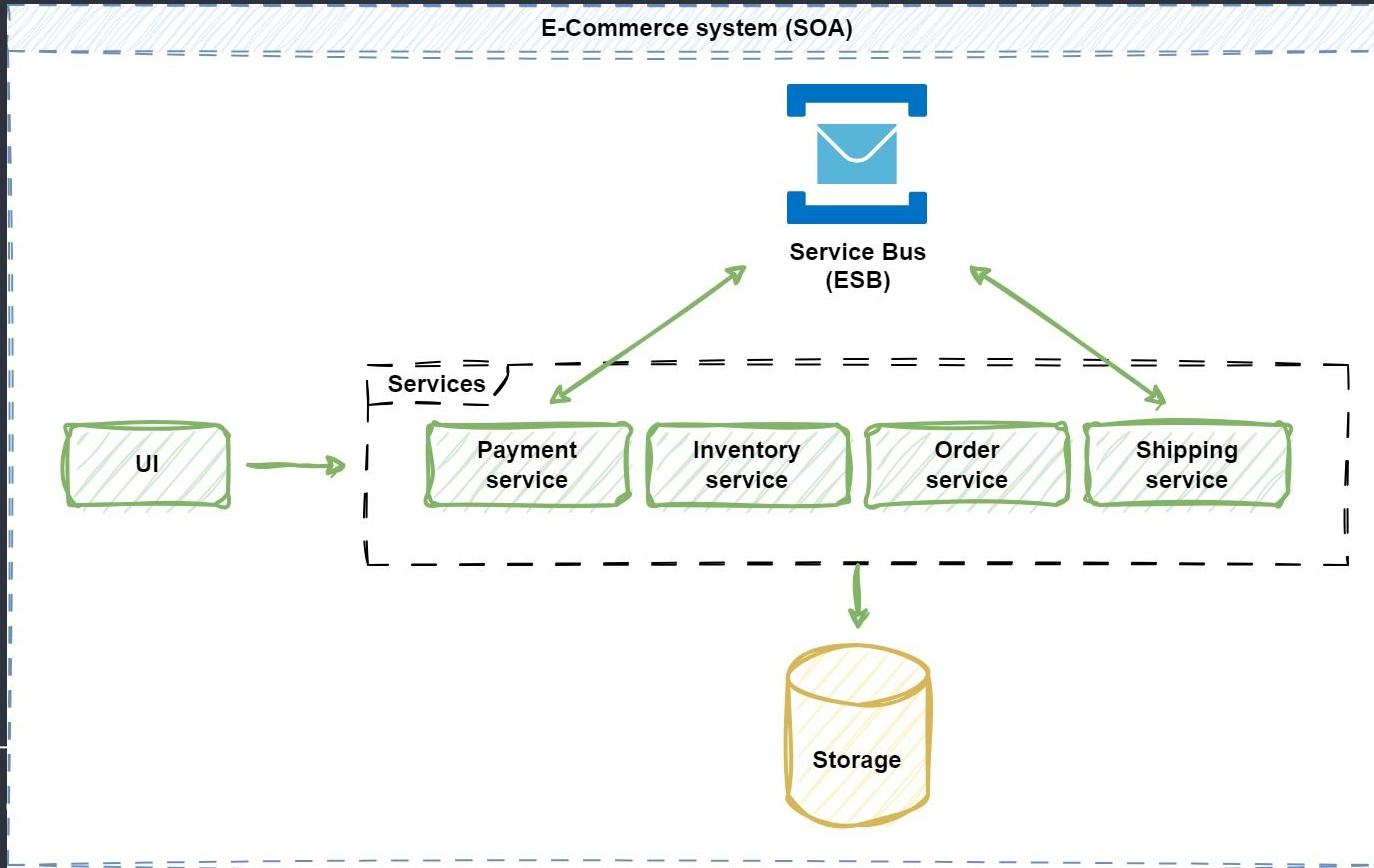
- A collection of independent components located on different machines that share messages with each other in order to achieve common goals
- The most important functions of distributed system
 - + Resource sharing
 - + Openness
 - + Concurrency
 - + Scalability
 - + Fault tolerance
 - + Transparency

Reference: <https://www.confluent.io/learn/distributed-systems/>

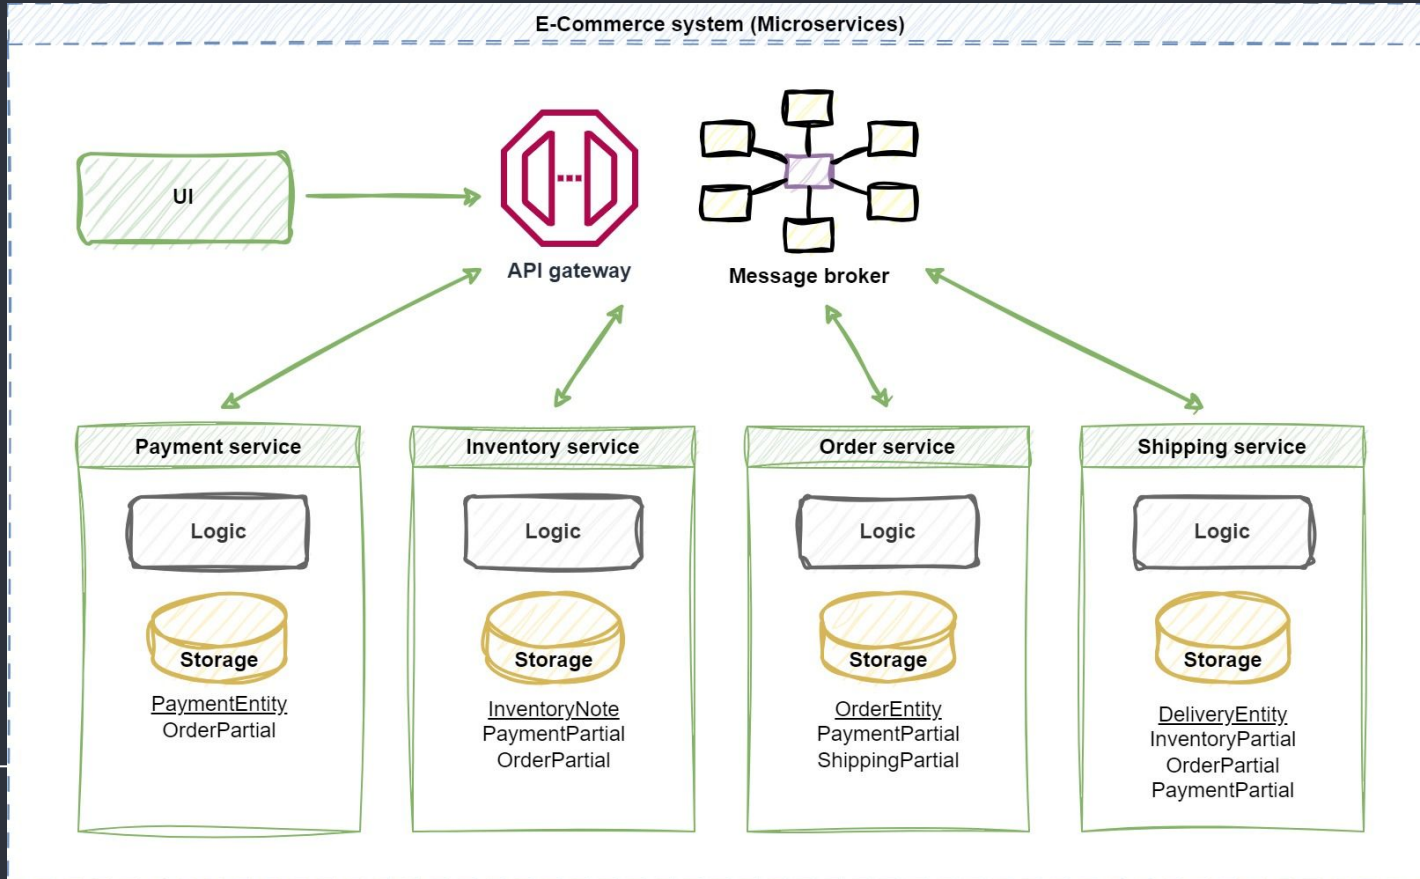
</ Monolithic - SOA - Microservices



</ Monolithic - SOA - Microservices




</ Monolithic - SOA - Microservices



</ Monolithic - SOA - Microservices

	Monolithic	SOA (Service-oriented arch)	Microservices
Scope	Enterprise	Enterprise	Application
Components	Single	Multiple	Multiple
Reuse	Code & logic	Components	Code & logic
Data duplication	Primary source	Primary source	Multiple sources
Communication	Synchronous	Mostly sync, sometimes async	Sync/Async
Channel	In-process	ESB (Enterprise Service Bus)	Message broker
Interoperability	In-process	Varied: REST, SOAP, AMQP, MSMQ	Varied (lightweight): REST, Protobuf, etc
Storage/Infras	Single	Single	Dedicated per service

</ Sample distributed order processing system



/>

</ Transaction in a distributed system

- **Atomicity:** a single unit with set of operations that must all occur or none occur
- **Consistency:** data should only move from one state to another expected valid state in all participants
- **Isolation:** ensures that concurrent transactions result in the same outcome when running sequentially
- **Durability:** makes the commit status of transactions persistent and be able to bear with failure related to system or power outage

</ Problems with a distributed transaction

Problems with transaction in a distributed environment


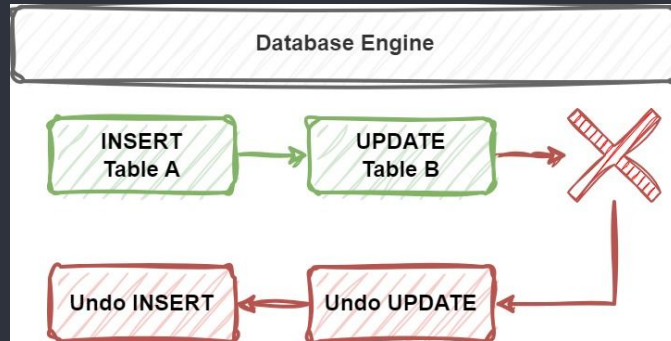
- Due to usage of database-per-microservice, NoSQL databases, message brokers, file storage
 - > They do not support **Two-phase commit (2PC)** protocol



</ Problems with a distributed transaction

Problems with transaction in a distributed environment

- Complicated rollback operations

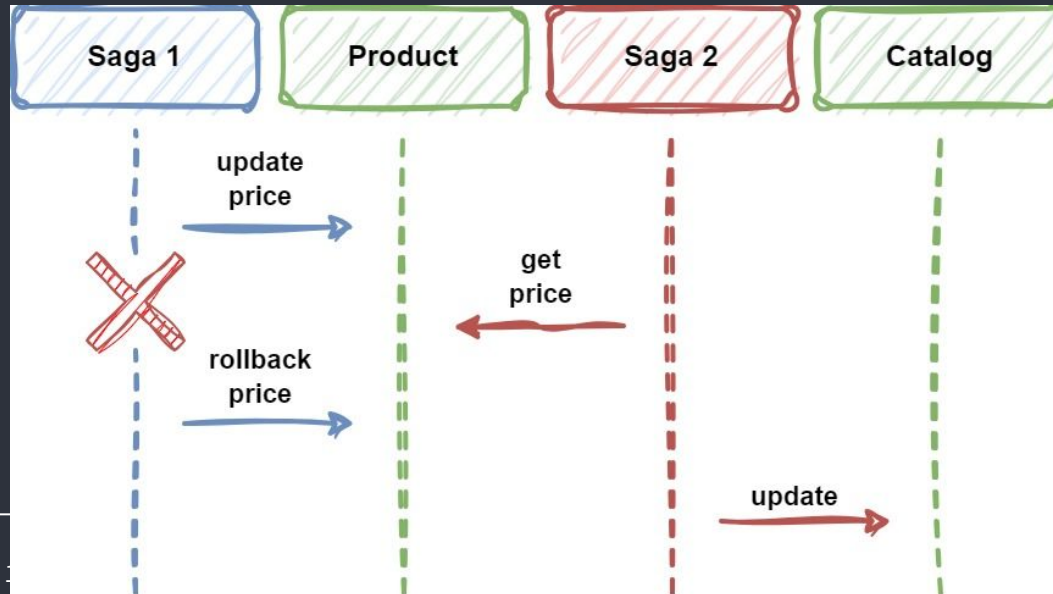


1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1

</ Problems with a distributed transaction

Problems with transaction in a distributed environment

- Atomicity, consistency and isolation is a challenge




/>

</ Problems with a distributed transaction

Problems with transaction in a distributed environment

- Limitations with synchronicity and availability



/>

</>

Saga - The solution

02

} /> [


1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1 0 1

</ What is Saga pattern?

- A transaction management strategy involving with multiple local transactions
- A local transaction is performed by a saga participant
- Each local transaction performs its transaction, then publishes events to trigger the next transaction
- If local transaction fails, the saga will execute a series of compensating transactions that reverse/compensate the changes made before



</ What is Saga pattern?



</ Saga pattern - Concepts

- Saga participants
- Compensable & compensating transactions
- Pivot transaction
- Retryable transactions
- Transaction key



</ Saga participants

Services/software components that are responsible for one or more local transactions in the whole saga transaction



</ Compensable, compensating transactions

- **Compensable transactions** can potentially be undo by executing another transaction with the opposite effect
- **Compensating transactions** reverse/compensate the changes made before by the compensable transactions



</ Pivot transaction

- The go/no-go (all or nothing) point in a saga
- If the pivot transaction commits, the saga runs the remaining transactions until done
- We can tell the saga is successful if a pivot transaction commits
- A pivot transaction can be neither compensable nor retryable, or it can be the last compensable transaction or the first retryable transaction in the saga



</ Retryable transactions

- Follow pivot transaction and are guaranteed to succeed
- For example, in ordering application
 - Create order
 - **Compensating transaction** will be Cancel order
 - Process payment
 - Update inventory
 - Delivery to customer → **pivot transaction**
 - Complete order → is **retryable** since all important steps are already successful

</ Transaction key

- Each transaction should have a unique and consistent key across services
- Is used to retrieve the transaction's information
- Marks entities or operations as associated with a specific transaction
- Can be randomly generated and linked or must be able to construct using consistent hashing



</ Implementation approaches

There are two common implementations:

- Choreography
- Orchestration

Each has its own set of challenges and techniques to manage the workflow



</>

Techniques in Saga

03

} /> [


1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 0 1

</ Common techniques & design patterns

- Asynchronous communication & Message broker
- Asynchronous request-reply
- "Replay" technique
- Other techniques to ensure resiliency and target isolation, consistency issues



</ Asynchronous communication & Message broker




</ Synchronous vs Asynchronous communication

- Asynchronous fits Saga better

	Synchronous	Asynchronous
Technique	HTTP request	Message queue/broker
Response time	Client waits for response	Quick response
Retry	Client waits for retry	Retry in background
Throughput	Lower	Higher
Timeout	Has timeout	Almost no timeout

</ Asynchronous request-reply



</ "Replay" technique

- A group of logical operations can be replayed multiple times
- Each operation is coordinated by the orchestrator
- Operation status and result will be persisted
- In the next replay, persisted result will be used if available
 - > Is usually used with Event-sourcing architecture
 - > More clear and readable but complicated



</ "Replay" technique

```
try
{
    ThrowIfFailed(transaction);

    OrderModel order = await GetCreatedOrder(transaction);

    // 1st play
    PaymentModel payment = await ProcessPayment(order, transaction, transactionService, publisher);

    // 2nd play
    InventoryNoteModel inventoryNote = await InventoryDelivery(payment, transaction, transactionService, publisher);

    // 3rd play
    DeliveryModel delivery = await ProcessDelivery(inventoryNote, transaction, transactionService, publisher);

    // 4th play
    await CompleteOrder(delivery, transaction, transactionService, publisher);
}
catch (TransactionFailedException ex)
{
    // x? play (with failure)
    await Rollback(ex, transaction, transactionService, publisher);
}
catch (AsyncTransactionException ex)
{
    Console.WriteLine(ex.Message);
}
```

</ Other techniques

To ensure resiliency and target isolation, consistency issues

- Heartbeat pattern
- Retry
- Semantic lock
- Optimistic locking
- Different concurrent read techniques



</>

Saga - Choreography

04

} /> [


1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1

</ What is choreography-based saga?

- A way to coordinate sagas where events are exchanged using a message broker or any mechanism without a centralized controller
- Each local transaction publishes events that trigger local transactions in other participants
- Participants must know each other's events to be able to handle them



</ What is choreography-based saga?



`</` Apply to our problem (Success path)

1. Create order -> ?

`/>`

</ Apply to our problem (Success path)

1. Create order -> Order created ->
Process payment -> Payment created ->
Inventory delivery -> Delivery note created ->
Process delivery -> Delivery created ->
Complete order -> Order completed



</ Apply to our problem (Failure paths)

1. Create order -> Create order failed
2. ... Process payment -> Payment failed ->
Cancel order
3. ... Inventory delivery -> Inventory delivery failed ->
Cancel order, Cancel payment
4. ... Process delivery -> Delivery failed ->
Cancel order, Cancel payment, Reverse inventory delivery

</ Apply to our demo [Publishing events]

```
public async Task PublishOrderCreated(OrderModel model)
{
    string message = JsonConvert.SerializeObject(new OrderCreatedEvent
    {
        Model = model
    });

    await _producer.ProduceAsync(nameof(OrderCreatedEvent),
        new Message<string, string>
        {
            Key = model.Id.ToString(),
            Value = message
        });
}
```

OrderService



</ Apply to our demo [Handling events]

2 references

```
public async Task HandleCreatePaymentWhenOrderCreated(CancellationTokentoken cancellationToken)
{
    await StartConsumerThread(nameof(OrderCreatedEvent), async (message) =>
    {
        using IServiceScope scope = serviceProvider.CreateScope();

        OrderCreatedEvent @event = JsonConvert.DeserializeObject<OrderCreatedEvent>(message.Message.Value);

        IPaymentService paymentService = scope.ServiceProvider.GetRequiredService<IPaymentService>();

        await paymentService.CreatePaymentFromOrder(@event.Model);
    }, cancellationToken: cancellationToken);
}
```

PaymentService



</ Choreography - Pros & Cons

Pros

- Good for simple, few participants, no need coordination logic workflows
- No need to maintain orchestrator service
- No single point of failure
- Responsibilities are shared and distributed

Cons

- Cumbersome when adding new participants and steps
- Complicated dependencies (cyclic, tight coupling) between participants
- All services must be running to simulate a transaction -> difficult to do integration testing

</>

Saga - Orchestrator

05

} /> [


1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1 0 1

</ What is orchestrator-based saga?

- Orchestrator coordinates sagas using a centralized controller
- Trigger participants local transactions to perform in case of events
- Handle failures by running compensating transactions/actions
- Status can be managed and persisted for monitoring and management purposes



</ What is orchestrator-based saga?



>

</ Apply to our problem (Success path)

1. Create order -> Order created -> ?

/>

</ Apply to our problem (Success path)

1. Create order -> Order created -> Orc ->
C:Process payment -> Payment created -> Orc ->
C:Inventory delivery -> Delivery note created -> Orc ->
C:Process delivery -> Delivery created -> Orc ->
C:Complete order -> Order completed -> Orc

Annotation

Orc: Orchestrator

C:{action}: command that triggers action

</ Apply to our problem (Failure paths)

1. ... Process payment -> **Payment failed** -> Orc ->
C:Cancel order
2. ... Inventory delivery -> **Inventory ... failed** -> Orc ->
C:Cancel order, C:Cancel payment
3. ...

Annotation

Orc: Orchestrator

C:{action}: command that triggers action

</ Apply to our demo [Orchestrator handles events]

1 reference

```
protected async Task HandleOrderCreated(Cancellation_token cancellationToken)
{
    await StartConsumerThread(nameof(OrderCreatedEvent), async (message) =>
    {
        using IServiceScope scope = serviceProvider.CreateScope();

        IOrderProcessingPublisher publisher = scope.ServiceProvider.GetRequiredService<IOrderProcessingPublisher>();

        OrderCreatedEvent @event = JsonConvert.DeserializeObject<OrderCreatedEvent>(message.Message.Value);

        await publisher.ProcessPayment(@event.Model);
    }, cancellationToken: cancellationToken);
}
```

Orchestrator

</ Apply to our demo (Orchestrator publishes commands)

3 references

```
public async Task ProcessPayment(OrderModel order)
{
    string message = JsonConvert.SerializeObject(new ProcessPaymentCommand
    {
        FromOrder = order,
    });

    await _producer.ProduceAsync(nameof(ProcessPaymentCommand),
        new Message<string, string>
        {
            Key = order.Id.ToString(),
            Value = message
        });
}
```

Orchestrator

</ Apply to our demo [Services handle commands]

2 references

```
public async Task HandleProcessPayment(CancellationTokentoken cancellationTokentoken)
{
    await StartConsumerThread(nameof(ProcessPaymentCommand), async (message) =>
    {
        using IServiceScope scope = serviceProvider.CreateScope();

        ProcessPaymentCommand command = JsonConvert.DeserializeObject<ProcessPaymentCommand>(message.Message.Value);

        IPaymentService paymentService = scope.ServiceProvider.GetRequiredService<IPaymentService>();

        await paymentService.CreatePaymentFromOrder(command.FromOrder);
    }, cancellationTokentoken: cancellationTokentoken);
}
```

PaymentService

</ Orchestrator - Pros & Cons

Pros

- Suitable for complex, many participants or new participants added over time workflows
- Good when flow of tasks and activities are under control
- No cyclic dependencies
- Separation of concerns

Cons

- Additional effort to maintain, implement coordination logic
- Single point of failure



</>

Issues & considerations

06

} /> [

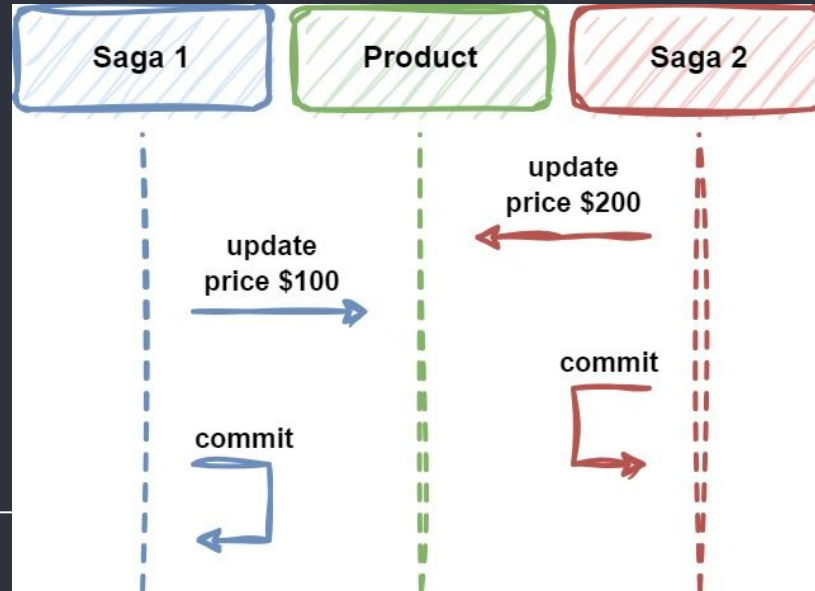
1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1 0 1

</ Some considerations

- Can be challenging in the beginning
- Hard to debug
- Complexity grows when there are more participants
- Consistency is relative only, because changes are committed to local databases
- Improper transient failures handling can cause side-effects
- Difficult to manage without proper monitoring and observability implemented

</ Common issues - Lost updates


- Lost updates, a participant updates data without reading latest changes



>

</ Common issues - Dirty reads


- Dirty reads, when a participant reads data that is just partially updated by another saga



/>

</ Common issues - Fuzzy/non-repeatable reads

- Fuzzy/non-repeatable reads, data is not consistent across participants since data can be updated at some points throughout the entire saga workflow



/>

</ Countermeasures against isolation issues

- Semantic lock
- Commutative updates
- Pessimistic view
- Reread value and restart if data changed
- Record operations in arriving order to execute sequentially
- Low-risk requests use sagas, high-risk requests favor distributed transactions (2PC, 3PC distributed transaction commit/rollback protocol)

</ Solutions for resiliency

- Rate limiting
- Delay & retry
- Idempotency



</>

Practical usage

07

} /> [

1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1 0 1

</ When to use

- Ensure data consistency in a distributed system without tight coupling
- Efficiently execute compensating actions if one of the operations fails in the distributed action sequence



</ When to use

- Microservices
- Cloud-based applications
 - [Saga pattern - AWS Prescriptive Guidance \(amazon.com\)](#)
 - [Saga pattern - Azure Design Patterns | Microsoft Learn](#)
 - [Saga pattern - IBM Cloud Architecture Center](#)



</ When to use

- Banking applications: [Patterns and implementations for a banking cloud transformation - Azure Architecture Center | Microsoft Learn](#)
- Media processing system: [Gridwich saga orchestration - Azure Reference Architectures | Microsoft Learn](#)



</ When to NOT use

- Monolithic architecture
- Tightly coupled operations and transactions
- Cyclic dependencies



</ Summary

- So far, we have learned that:
 - The use of Saga pattern in distributed transaction
 - Some common techniques used in Saga
 - Two Saga implementations: Choreography & Orchestrator
 - Some common issues and challenging when apply Saga



</>

Discussion & QnA

08

} /> [

1 0 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 0 1 1 1 0 1 1 1 0 1

</ References

- [SOA vs. Microservices: What's the Difference? | IBM](#)
- [Saga pattern - Azure Design Patterns | Microsoft Learn](#)
- [Patterns and implementations for a banking cloud transformation - Azure Architecture Center | Microsoft Learn](#)
- [Gridwich saga orchestration - Azure Reference Architectures | Microsoft Learn](#)
- [Saga Pattern in Microservices | Baelung on Computer Science](#)
- [Saga Pattern for Microservices Distributed Transactions | by Mehmet Ozkaya | Design Microservices Architecture with Patterns & Principles | Medium](#)
- [Distributed transaction - SAGA pattern \(viblo.asia\)](#)
- [Distributed Systems: An Introduction to Distributed Computing \(confluent.io\)](#)



</ Thanks!

Do you have any questions?

trannamtrung1st@gmail.com

<https://www.linkedin.com/in/trung-tran99/>



/>

} /> [

CREDITS: This presentation template was created by Slidesgo, and includes icons by Flaticon, and infographics & images by Freepik

Please keep this slide for attribution